

Lewis Arnsten

I apologize for this homeworks slight lateness. I had COVID up until the middle of this week (I tested positive Feb. 22). This significantly impaired my ability to do work over the course of this assignment.

### Image of gradient\_checker results:

```
[(base) Lewiss-MacBook-Pro:hw3 lewisarnsten$ python3 gradient_checker.py  
Testing Conv2d:  
Shape Match = True  
Result Match = True  
Result Match = False  
Result Match = True  
Result Match = True  
  
Testing Linear:  
Shape Match = True  
Result Match = True  
Result Match = True  
Result Match = True  
Result Match = True  
  
Testing AvgPool2d:  
Shape Match = True  
Result Match = True  
Result Match = True  
  
Testing ELU:  
Shape Match = True  
Result Match = True  
Result Match = True  
  
Testing BatchNorm2d:  
Shape Match = True  
Result Match = True  
Result Match = True  
Result Match = True  
Result Match = True  
  
Testing cross_entropy_loss_with_softmax:  
Shape Match = True  
Result Match = True  
Result Match = True
```

### Logic behind pytorch CNN architecture for segmentation:

I first created a series of sequences (4 sequences) to conduct convolutions, batch normalization, and ReLU activation--adding layers for feature extraction. When each sequence was called, it was followed by max pooling with stride two to cut image resolution in half. In my first convolution I increased the number of layers from 3 to 64. I then followed a common practice of doubling the number of layers in each convolution until I felt it would be trivial to add more (512). I then created and executed 4 more sequences utilizing convolutions, batch normalization, ReLU activation, and transpose convolution resulting in output of shape (64, 20, 32, 32). I tested the SGD and Adam optimizers, and decided upon Adam due to better performance. I also decided upon a learning rate of  $3e-4$  experimentally.

**Results: achieved ~80% accuracy in 20 epochs**

```
Epoch 18, Iteration 0, loss = 0.1236
Epoch 18, Iteration 100, loss = 0.1896
Epoch 18, Iteration 200, loss = 0.1141
Epoch 18, Iteration 300, loss = 0.1157
Epoch 18, Iteration 400, loss = 0.1532
Epoch 18, Iteration 500, loss = 0.1926
Epoch 18, Iteration 600, loss = 0.3634
Epoch 18, Iteration 700, loss = 0.2351
Epoch 19, Iteration 0, loss = 0.1206
Epoch 19, Iteration 100, loss = 0.1905
Epoch 19, Iteration 200, loss = 0.0894
Epoch 19, Iteration 300, loss = 0.0819
Epoch 19, Iteration 400, loss = 0.1238
Epoch 19, Iteration 500, loss = 0.1598
Epoch 19, Iteration 600, loss = 0.1443
Epoch 19, Iteration 700, loss = 0.0831
Got 8165912 / 10240000 correct (79.75)
```